

Title TBD

Rodrigo Araújo

University of British Columbia
Vancouver, British Columbia
rodarauj@cs.ubc.ca

Reid Holmes

University of British Columbia
Vancouver, British Columbia
rtholmes@cs.ubc.ca

ABSTRACT

Due to the advancement of distributed systems and the increasing industrial demands, software systems contain multiple components with complex interactions, e.g databases and their replication, caching components, proxies and load balancers, application instances and their complex configuration parameters. The engineers in a project must think with many configuration parameters that change the behavior and/or structure of the system, this can cause many problems that affect the quality of the service. In other words, dealing with high dimensionality is both cognitively demanding and risky for the project.

In this work we show the design and analysis of a pragmatic machine learning based tool that aims to assist the engineering of systems that can: 1) monitor themselves, 2) Forecast workloads and performance metrics and 3) Change themselves in run-time by self-configuring and adapting for a specific scenario. After the integration of this tool with a system, it should be able to answer the question: given that we have many configuration parameters, how can we change them in order to optimize a certain metric for a given predicted workload?

We show that it can decrease the risk of changing systems' configurations in run-time and decrease the engineering effort that otherwise would be spent manually optimizing parameters, usually following a trial-and-error approach.

CCS CONCEPTS

• TBD → TBD;

KEYWORDS

ACM proceedings, L^AT_EX, text tagging

ACM Reference Format:

Rodrigo Araújo and Reid Holmes. 1997. Title TBD. In *Proceedings of ACM Woodstock conference (WOODSTOCK'97)*, Rodrigo Araújo and Reid Holmes (Eds.). ACM, New York, NY, USA, Article 4, 2 pages. https://doi.org/10.475/123_4

1 INTRODUCTION

An extended version of the abstract, plus some background information, introduction to the problem being solved then a discussion about the approach to solve the problem. Brief discussion of evaluation.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

WOODSTOCK'97, July 1997, El Paso, Texas USA

© 2016 Copyright held by the owner/author(s).

ACM ISBN 123-4567-24-567/08/06...\$15.00

https://doi.org/10.475/123_4

1.1 Motivation

Here I should try to show what is the problem by giving a motivation. Followed by showing why this problem is both interesting and important and why do naive approaches usually fail. Followed by why it hasn't been solved before

1.2 Definition of the problem

Here I try to better define the problem in a formal manner.

1.3 Proposed approach

Here I should give the key components of my approach and results, followed by the structure of the paper.

2 RELATED WORK

Maybe related work would be better now instead of putting it in the end, since there are a lot of related work and they are very relevant to this work, it would be nice to upfront point out what has been done and what would be different in this work.

3 APPROACH

3.1 Control theory and self-adaptive systems

3.2 System's configuration as an optimization problem

3.3 Providing system adaptation with machine learning

3.4 Workload simulation

3.5 System instrumentation

3.6 Machine learning architecture

3.6.1 Features and models.

3.6.2 Online training.

3.6.3 Achieving self-adaptation.

4 IMPLEMENTATION

5 EVALUATION AND DISCUSSION

6 FUTURE WORK

7 CONCLUSIONS

ACKNOWLEDGMENTS

REFERENCES

- [1] [n. d.]. short term performance forecasting in enterprise systems. ([n. d.]). <http://www.hpl.hp.com/techreports/2005/HPL-2005-50.pdf>
- [2] [n. d.]. Using probabilistic reasoning to automate software tuning. ([n. d.]). <http://ftp.deas.harvard.edu/techreports/tr-08-04.pdf>

- [3] Virgilio AF Almeida. 2002. Capacity Planning for Web Services.. In *Performance*. Springer, 142–157. <http://link.springer.com/content/pdf/10.1007/3-540-45798-4.pdf#page=151>
- [4] Joy Arulraj Haibin Lin Jiexi Lin Lin Ma Prashanth Menon Todd Mowry Matthew Perron Ian Quah Siddharth Santurkar Anthony Tomasic Skye Toor Dana Van Aken Ziqi Wang Yingjun Wu Ran Xian Tieying Zhang Andrew Pavlo, Gustavo Angulo. 2017. *Self-Driving Database Management Systems*.
- [5] Ira Cohen, Jeffrey S. Chase, Moises Goldszmidt, Terence Kelly, and Julie Symons. 2004. Correlating Instrumentation Data to System States: A Building Block for Automated Diagnosis and Control.. In *OSDI*, Vol. 4. 16–16. <http://sailhome.cs.queensu.ca/~corpaul/ciscxxx/papers/HPL-2004-183.pdf>
- [6] Rogério de Lemos, Holger Giese, Hausi A. Müller, and Mary Shaw (Eds.). 2013. *Software Engineering for Self-Adaptive Systems II*. Lecture Notes in Computer Science, Vol. 7475. Springer Berlin Heidelberg, Berlin, Heidelberg. <http://link.springer.com/10.1007/978-3-642-35813-5> DOI: 10.1007/978-3-642-35813-5.
- [7] Serge Demeyer. 2011. Research methods in computer science.. In *ICSM*. 600. http://ansymore.uantwerpen.be/system/files/uploads/courses/CapitaSelecta/ResearchMethods00_Contents.pdf
- [8] Dirk Draheim, John Grundy, John Hosking, Christof Lutteroth, and Gerald Weber. 2006. Realistic load testing of web applications. In *Software Maintenance and Reengineering, 2006. CSMR 2006. Proceedings of the 10th European Conference on*. IEEE, 11–pp. <http://ieeexplore.ieee.org/abstract/document/1602358/>
- [9] F. Faniyi, P. R. Lewis, R. Bahsoon, and X. Yao. 2014. Architecting Self-Aware Software Systems. In *2014 IEEE/IFIP Conference on Software Architecture*. 91–94. <https://doi.org/10.1109/WICSA.2014.18>
- [10] Archana Sulochana Ganapathi. 2009. *Predicting and Optimizing System Utilization and Performance via Statistical Machine Learning*. Ph.D. Dissertation. EECS Department, University of California, Berkeley. <http://www2.eecs.berkeley.edu/Pubs/TechRpts/2009/EECS-2009-181.html>
- [11] Martin Georgiev, Subodh Iyengar, Suman Jana, Rishita Anubhai, Dan Boneh, and Vitaly Shmatikov. 2012. The most dangerous code in the world: validating SSL certificates in non-browser software. In *Proceedings of the 2012 ACM conference on Computer and communications security*. ACM, 38–49. <http://dl.acm.org/citation.cfm?id=2382204>
- [12] Nikolas Roman Herbst, Nikolaus Huber, Samuel Kounev, and Erich Amrehn. 2014. Self-adaptive workload classification and forecasting for proactive resource provisioning. *Concurrency and Computation: Practice and Experience* 26, 12 (Aug. 2014), 2053–2078. <https://doi.org/10.1002/cpe.3224>
- [13] Shijun Liu, Zeyu Di, Lei Wu, Li Pan, and Yuliang Shi. 2016. Probabilistic-based workload forecasting and service redeployment for multi-tenant services. *International Journal of High Performance Computing and Networking* 9, 1/2 (2016), 134. <https://doi.org/10.1504/IJHPCN.2016.074666>
- [14] Theophano Mitsa. 2010. *Temporal data mining*. Chapman & Hall/CRC, Boca Raton, FL. OCLC: ocn474869534.
- [15] Jeho Oh, Don Batory, Margaret Myers, and Norbert Siegmund. 2017. Finding near-optimal configurations in product lines by random sampling. ACM Press, 61–71. <https://doi.org/10.1145/3106237.3106273>
- [16] Barry Porter, Matthew Grieves, Roberto Rodrigues Filho, and David Leslie. 2016. REX: A Development Platform and Online Learning Approach for Runtime Emergent Software Systems. In *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*. USENIX Association, GA, 333–348. <https://www.usenix.org/conference/osdi16/technical-sessions/presentation/porter>
- [17] M. Salehie and L. Tahvildari. 2009. Self-adaptive software: Landscape and research challenges. *ACM Transactions on Autonomous and Adaptive Systems (TAAS)*.
- [18] Stepan Shevtsov, Mihaly Berekmeri, Danny Weyns, and Martina Maggio. 2017. Control-Theoretical Software Adaptation: A Systematic Literature Review. *IEEE Transactions on Software Engineering* (2017), 1–1. <https://doi.org/10.1109/TSE.2017.2704579>